

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPClassifier
!pip install dmba
from dmba import classificationSummary
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import r2_score, mean_squared_error
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeClassifier
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import precision_score, recall_score, f1_score

```

[Show hidden output](#)

Upload files from local system

```

from google.colab import files
up = files.upload()

Choose Files loan_approval.csv
loan_approval.csv(text/csv) - 115793 bytes, last modified: 11/7/2025 - 100% done
Saving loan_approval.csv to loan_approval (1).csv

```

Read the loan approval dataset

```

loan_df = pd.read_csv('loan_approval.csv')
print(loan_df.head(5))

      name        city   income  credit_score  loan_amount \
0  Allison Hill    East Jill  113810          389       39698
1  Brandon Hall  New Jamesside  44592          729       15446
2  Rhonda Smith    Lake Roberto  33278          584       11189
3 Gabrielle Davis  West Melanieview 127196          344       48823
4  Valerie Gray    Mariastad  66048          496       47174

  years_employed  points  loan_approved
0            27    50.0        False
1            28    55.0        False
2            13    45.0        False
3            29    50.0        False
4             4    25.0        False

```

Check for missing values and information

```

print(loan_df.isnull().sum())
loan_df.info()

name      0
city      0
income     0
credit_score     0
loan_amount     0
years_employed     0
points      0
loan_approved     0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   name        2000 non-null    object 
 1   city         2000 non-null    object 
 2   income       2000 non-null    int64  
 3   credit_score 2000 non-null    int64  
 4   loan_amount  2000 non-null    int64  

```

```

5    years_employed  2000 non-null  int64
6    points          2000 non-null  float64
7    loan_approved   2000 non-null  bool
dtypes: bool(1), float64(1), int64(4), object(2)
memory usage: 111.5+ KB

```

▼ Create copies and drop unnecessary columns

```

loan_df2 = loan_df
loan_df3 = loan_df
loan_df4 = loan_df
loan_df = loan_df.drop(['name', 'city', 'points'], axis=1)
print(loan_df.head())

income  credit_score  loan_amount  years_employed  loan_approved
0    113810           389        39698            27      False
1     44592            729        15446            28      False
2    33278             584        11189            13      False
3    127196             344        48823            29      False
4    66048             496        47174             4      False

```

▼ NN classification:

▼ Define predictor and outcome for NN classification

```

predictors = ['income', 'credit_score', 'loan_amount', 'years_employed']
outcome = 'loan_approved'

predictors = [c for c in loan_df.columns if c != outcome]
print(loan_df.shape)

(2000, 5)

```

▼ Split data into training data and validation data, and scale features

```

X = loan_df[predictors]
y = loan_df[outcome]
train_X, valid_X, train_y, valid_y = train_test_split(X, y, test_size=0.4, random_state=1)

scaler = StandardScaler()
scaler.fit(train_X)

train_X = scaler.transform(train_X)
valid_X = scaler.transform(valid_X)

```

▼ Train an NN classifier

```

clf = MLPClassifier(hidden_layer_sizes=(6), activation='logistic', solver='adam',
random_state=1, max_iter=10000)
clf.fit(train_X, train_y.values)

```

MLPClassifier	?
MLPClassifier(activation='logistic', hidden_layer_sizes=6, max_iter=10000, random_state=1)	

▼ classification summary

```

classificationSummary(train_y, clf.predict(train_X))
classificationSummary(valid_y, clf.predict(valid_X))

```

Confusion Matrix (Accuracy 0.9075)

Prediction

```
Actual   0   1
      0  614  55
      1  56 475
Confusion Matrix (Accuracy 0.8975)
```

```
Prediction
Actual   0   1
      0  412  40
      1  42 306
```

```
valid_pred = clf.predict(valid_X)
precision = precision_score(valid_y, valid_pred)
recall = recall_score(valid_y, valid_pred)
f1 = f1_score(valid_y, valid_pred)

print(f"NN Classifier Metrics:")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")
```

```
NN Classifier Metrics:
Precision: 0.8844
Recall: 0.8793
F1-Score: 0.8818
```

▼ NN regression:

▼ Drop unnecessary columns

```
loan_df2 = loan_df2.drop(['name', 'city', 'loan_approved'], axis=1)
print(loan_df2.head())
```

	income	credit_score	loan_amount	years_employed	points
0	113810	389	39698	27	50.0
1	44592	729	15446	28	55.0
2	33278	584	11189	13	45.0
3	127196	344	48823	29	50.0
4	66048	496	47174	4	25.0

▼ Define predictors and outcome for NN regression

```
predictors2 = ['income', 'credit_score', 'loan_amount', 'years_employed']
outcome2 = 'points'
```

▼ Split data into training data and validation data, and scale features

```
X2 = loan_df2[predictors2]
y2 = loan_df2[outcome2]
train_X2, valid_X2, train_y2, valid_y2 = train_test_split(X2, y2, test_size=0.4, random_state=1)

scaler2 = StandardScaler()
scaler2.fit(train_X2)

train_X2 = scaler2.transform(train_X2)
valid_X2 = scaler2.transform(valid_X2)
```

▼ Train an NN regressor

```
clf2 = MLPRegressor(hidden_layer_sizes=(6), activation='relu', solver='adam',
random_state=1, max_iter=10000)
clf2.fit(train_X2, train_y2.values)
```

```
MLPRegressor(hidden_layer_sizes=6, max_iter=10000, random_state=1)
```

▼ NN regression model performance

```

train_pred = clf2.predict(train_X2)
valid_pred = clf2.predict(valid_X2)

train_r2 = r2_score(train_y2, train_pred)
valid_r2 = r2_score(valid_y2, valid_pred)
train_rmse = np.sqrt(mean_squared_error(train_y2, train_pred))
valid_rmse = np.sqrt(mean_squared_error(valid_y2, valid_pred))

print("Train:")
print("R-squared:", train_r2)
print("RMSE:", train_rmse)

print("Valid:")
print("R-squared:", valid_r2)
print("RMSE:", valid_rmse)

Train:
R-squared: 0.9295515300418117
RMSE: 4.98450161674012
Valid:
R-squared: 0.923129384117913
RMSE: 5.10474437354557

```

▼ NN regression Cross-Validation

```

scores_clf2 = cross_val_score(
    clf2, X2, y2, cv=5,
    scoring='neg_root_mean_squared_error')

scores_clf2 = -scores_clf2

print("NN regression Cross-Validation:")
print(f"RMSE Scores (5-fold): {np.round(scores_clf2, 3)}")
print(f"Mean RMSE: {scores_clf2.mean():.3f}")
print(f"RMSE Std Dev: {scores_clf2.std():.3f}")

NN regression Cross-Validation:
RMSE Scores (5-fold): [14.205 12.755 13.41  12.768 13.291]
Mean RMSE: 13.286
RMSE Std Dev: 0.531

```

▼ CART regression:

▼ Define predictor and outcome for CART

```

predictors3 = ['income', 'credit_score', 'loan_amount', 'years_employed']
outcome3 = 'points'

X3 = loan_df3[predictors3]
y3 = loan_df3[outcome3]
train_X3, valid_X3, train_y3, valid_y3 = train_test_split(X3, y3, test_size=0.4, random_state=1)

```

▼ Train an CART regressor with max_depth=8

```

cart_r8 = DecisionTreeRegressor(max_depth=8, random_state=1)
cart_r8.fit(train_X3, train_y3.values)

train_pred_cart_r8 = cart_r8.predict(train_X3)
valid_pred_cart_r8 = cart_r8.predict(valid_X3)

cart_r8_train_r3 = r2_score(train_y3, train_pred_cart_r8)
cart_r8_valid_r3 = r2_score(valid_y3, valid_pred_cart_r8)
cart_r8_train_rmse3 = np.sqrt(mean_squared_error(train_y3, train_pred_cart_r8))
cart_r8_valid_rmse3 = np.sqrt(mean_squared_error(valid_y3, valid_pred_cart_r8))

```

▼ CART8 regression model performance

```

print("CART_r8 Train:")
print("R-squared:", cart_r8_train_r3)
print("RMSE:", cart_r8_train_rmse3)

print("CART_r8 Valid:")
print("R-squared:", cart_r8_valid_r3)
print("RMSE:", cart_r8_valid_rmse3)

CART_r8 Train:
R-squared: 0.993666627854971
RMSE: 1.4945259573300949
CART_r8 Valid:
R-squared: 0.9653461079018368
RMSE: 3.427439742339933

```

▼ Train an CART regressor with max_depth=7

```

cart_r7 = DecisionTreeRegressor(max_depth=7, random_state=1)
cart_r7.fit(train_X3, train_y3.values)

train_pred_cart_r7 = cart_r7.predict(train_X3)
valid_pred_cart_r7 = cart_r7.predict(valid_X3)

cart_r7_train_r3 = r2_score(train_y3, train_pred_cart_r7)
cart_r7_valid_r3 = r2_score(valid_y3, valid_pred_cart_r7)
cart_r7_train_rmse3 = np.sqrt(mean_squared_error(train_y3, train_pred_cart_r7))
cart_r7_valid_rmse3 = np.sqrt(mean_squared_error(valid_y3, valid_pred_cart_r7))

```

▼ CART7 regression model performance

```

print("CART_r7 Train:")
print("R-squared:", cart_r7_train_r3)
print("RMSE:", cart_r7_train_rmse3)

print("CART_r7 Valid:")
print("R-squared:", cart_r7_valid_r3)
print("RMSE:", cart_r7_valid_rmse3)

CART_r7 Train:
R-squared: 0.9855979145954754
RMSE: 2.253713206104595
CART_r7 Valid:
R-squared: 0.9607408935212078
RMSE: 3.6480772344659123

```

▼ Train an CART regressor with max_depth=9

```

cart_r9 = DecisionTreeRegressor(max_depth=9, random_state=1)
cart_r9.fit(train_X3, train_y3.values)

train_pred_cart_r9 = cart_r9.predict(train_X3)
valid_pred_cart_r9 = cart_r9.predict(valid_X3)

cart_r9_train_r3 = r2_score(train_y3, train_pred_cart_r9)
cart_r9_valid_r3 = r2_score(valid_y3, valid_pred_cart_r9)
cart_r9_train_rmse3 = np.sqrt(mean_squared_error(train_y3, train_pred_cart_r9))
cart_r9_valid_rmse3 = np.sqrt(mean_squared_error(valid_y3, valid_pred_cart_r9))

```

▼ CART9 regression model performance

```

print("CART_r9 Train:")
print("R-squared:", cart_r9_train_r3)
print("RMSE:", cart_r9_train_rmse3)

print("CART_r9 Valid:")
print("R-squared:", cart_r9_valid_r3)
print("RMSE:", cart_r9_valid_rmse3)

CART_r9 Train:
R-squared: 0.9976059462347057
RMSE: 0.9188675366951881
CART_r9 Valid:
R-squared: 0.9674563081839435
RMSE: 3.3214462448298687

```

▼ Train an CART regressor with max_depth=10

```

cart_r10 = DecisionTreeRegressor(max_depth=10, random_state=1)
cart_r10.fit(train_X3, train_y3.values)

train_pred_cart_r10 = cart_r10.predict(train_X3)
valid_pred_cart_r10 = cart_r10.predict(valid_X3)

cart_r10_train_r3 = r2_score(train_y3, train_pred_cart_r10)
cart_r10_valid_r3 = r2_score(valid_y3, valid_pred_cart_r10)
cart_r10_train_rmse3 = np.sqrt(mean_squared_error(train_y3, train_pred_cart_r10))
cart_r10_valid_rmse3 = np.sqrt(mean_squared_error(valid_y3, valid_pred_cart_r10))

```

▼ CART10 regression model performance

```

print("CART_r10 Train:")
print("R-squared:", cart_r10_train_r3)
print("RMSE:", cart_r10_train_rmse3)

print("CART_r10 Valid:")
print("R-squared:", cart_r10_valid_r3)
print("RMSE:", cart_r10_valid_rmse3)

CART_r10 Train:
R-squared: 0.9992613598778436
RMSE: 0.5103902926539864
CART_r10 Valid:
R-squared: 0.9664697571568879
RMSE: 3.3714146389925452

```

▼ Cart_r9 Cross-Validation

```

scores_cart_r9 = cross_val_score(
    cart_r9, X3, y3, cv=5,
    scoring='neg_root_mean_squared_error')

scores_cart_r9 = -scores_cart_r9

print("Cart_r9 Cross-Validation:")
print(f"RMSE Scores (5-fold): {np.round(scores_cart_r9, 3)}")
print(f"Mean RMSE: {scores_cart_r9.mean():.3f}")
print(f"RMSE Std Dev: {scores_cart_r9.std():.3f}")

Cart_r9 Cross-Validation:
RMSE Scores (5-fold): [2.811 3.077 3.308 3.527 3.06]
Mean RMSE: 3.156
RMSE Std Dev: 0.243

```

▼ CART classification:

```

loan_df4 = loan_df4.drop(['name', 'city', 'points'], axis=1)

X4 = loan_df4[predictors]

```

```
y4 = loan_df4[outcome]
train_X4, valid_X4, train_y4, valid_y4 = train_test_split(X4, y4, test_size=0.4, random_state=1)
```

▼ Train an CART classification with max_depth=9

```
cart_c9 = DecisionTreeClassifier(max_depth=9, random_state=1)
cart_c9.fit(train_X4, train_y4)

classificationSummary(train_y4, cart_c9.predict(train_X4))
classificationSummary(valid_y4, cart_c9.predict(valid_X4))

Confusion Matrix (Accuracy 0.9983)

      Prediction
Actual   0   1
  0 667   2
  1   0 531

Confusion Matrix (Accuracy 0.9725)

      Prediction
Actual   0   1
  0 446   6
  1  16 332
```

▼ Train an CART classification with max_depth=8

```
cart_c8 = DecisionTreeClassifier(max_depth=8, random_state=1)
cart_c8.fit(train_X4, train_y4)

classificationSummary(train_y4, cart_c8.predict(train_X4))
classificationSummary(valid_y4, cart_c8.predict(valid_X4))

Confusion Matrix (Accuracy 0.9967)

      Prediction
Actual   0   1
  0 665   4
  1   0 531

Confusion Matrix (Accuracy 0.9712)

      Prediction
Actual   0   1
  0 445   7
  1  16 332
```

▼ Train an CART classification with max_depth=10

```
cart_c10 = DecisionTreeClassifier(max_depth=10, random_state=1)
cart_c10.fit(train_X4, train_y4)

classificationSummary(train_y4, cart_c10.predict(train_X4))
classificationSummary(valid_y4, cart_c10.predict(valid_X4))

Confusion Matrix (Accuracy 0.9983)

      Prediction
Actual   0   1
  0 669   0
  1   2 529

Confusion Matrix (Accuracy 0.9725)

      Prediction
Actual   0   1
  0 446   6
  1  16 332
```

▼ Cart_c9 Cross-Validation

```
scores_cart_c9 = cross_val_score(
    cart_c9, X4, y4, cv=5,
    scoring='neg_root_mean_squared_error')
```

```
scores_cart_c9 = -scores_cart_c9

print("CART classification Cross-Validation:")
print(f"RMSE Scores (5-fold): {np.round(scores_cart_c9, 3)}")
print(f"Mean RMSE: {scores_cart_c9.mean():.3f}")
print(f"RMSE Std Dev: {scores_cart_c9.std():.3f}")

CART classification Cross-Validation:
RMSE Scores (5-fold): [0.158 0.173 0.158 0.18  0.166]
Mean RMSE: 0.167
RMSE Std Dev: 0.009
```

```
valid_pred4 = cart_c9.predict(valid_X4)
precision = precision_score(valid_y4, valid_pred4)
recall = recall_score(valid_y4, valid_pred4)
f1 = f1_score(valid_y4, valid_pred4)

print(f"CART Classifier Metrics:")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")

CART Classifier Metrics:
Precision: 0.9822
Recall: 0.9540
F1-Score: 0.9679
```


```
sns.countplot(x='loan_approved', data=loan_df, palette='Set2')
plt.title('Figure 1. Distribution of Loan Approval Status')
plt.xlabel('Loan Approved (1 = Yes, 0 = No)')
plt.ylabel('Count')
plt.show()
```

/tmp/ipython-input-876353508.py:1: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set

```
sns.countplot(x='loan_approved', data=loan_df, palette='Set2')
```

Figure 1. Distribution of Loan Approval Status



```
plt.figure(figsize=(8,5))
sns.heatmap(loan_df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Figure 2. Feature Correlation Heatmap')
plt.show()
```



```

features = ['income', 'credit_score', 'loan_amount', 'years_employed']

for col in features:
    plt.figure(figsize=(5,3))
    sns.boxplot(x='loan_approved', y=col, data=loan_df, palette='Set1')
    plt.title(f'Figure: {col.capitalize()} vs Loan Approval')
    plt.xlabel('Loan Approved')
    plt.ylabel(col.capitalize())
    plt.show()


```

```
/tmp/ipython-input-1790707769.py:5: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set
```

```
sns.boxplot(x='loan_approved', y=col, data=loan_df, palette='Set1')
```

Figure: Income vs Loan Approval




```
/tmp/ipython-input-1790707769.py:5: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set
```

```
sns.boxplot(x='loan_approved', y=col, data=loan_df, palette='Set1')
```

Figure: Credit_score vs Loan Approval



```
/tmp/ipython-input-1790707769.py:5: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set
```

```
sns.boxplot(x='loan_approved', y=col, data=loan_df, palette='Set1')
```

Figure: Loan_amount vs Loan Approval

